

Weekly Assignment Report

Training Code:

Python

```
import torch

import torch.nn as nn

import torch.optim as optim

import torchvision

import torchvision.transforms as transforms

from torch.utils.data import DataLoader

from torch.optim.lr_scheduler import StepLR


# Data augmentation and normalization

transform_train = transforms.Compose([

    transforms.RandomCrop(32, padding=4),

    transforms.RandomHorizontalFlip(),

    transforms.RandomRotation(15),

    transforms.ToTensor(),

    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))

])


transform_test = transforms.Compose([

    transforms.ToTensor(),
```

```

        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))

    ])

# Load CIFAR-10 dataset

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
download=True, transform=transform_train)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
download=True, transform=transform_test)

trainloader = DataLoader(trainset, batch_size=256, shuffle=True, num_workers=2)

testloader = DataLoader(testset, batch_size=256, shuffle=False, num_workers=2)

# CNN model with normalization, pooling and convolution layers

class CNN_Model(nn.Module):

    def __init__(self):

        super(CNN_Model, self).__init__()

        self.conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)

        self.conv2 = nn.Conv2d(64, 128, kernel_size=3, padding=1)

        self.conv3 = nn.Conv2d(128, 256, kernel_size=3, padding=1)

        self.conv4 = nn.Conv2d(256, 512, kernel_size=3, padding=1)

        self.fc1 = nn.Linear(512 * 2 * 2, 1024)

        self.fc2 = nn.Linear(1024, 10)

        self.pool = nn.MaxPool2d(2, 2)

```

```

self.dropout = nn.Dropout(0.5)

self.batch_norm1 = nn.BatchNorm2d(64)

self.batch_norm2 = nn.BatchNorm2d(128)

self.batch_norm3 = nn.BatchNorm2d(256)

self.batch_norm4 = nn.BatchNorm2d(512)

def forward(self, x):

    x = self.pool(torch.relu(self.batch_norm1(self.conv1(x))))

    x = self.pool(torch.relu(self.batch_norm2(self.conv2(x))))

    x = self.pool(torch.relu(self.batch_norm3(self.conv3(x))))

    x = self.pool(torch.relu(self.batch_norm4(self.conv4(x))))

    x = x.view(-1, 512 * 2 * 2)

    x = torch.relu(self.fc1(x))

    x = self.dropout(x)

    x = self.fc2(x)

    return x

```

Initialize the model, loss function, optimizer, and learning rate scheduler

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
model = CNN_Model().to(device)
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
scheduler = StepLR(optimizer, step_size=30, gamma=0.1)
```

```
# Training loop

num_epochs = 30

for epoch in range(num_epochs):

    model.train()

    running_loss = 0.0

    correct = 0

    total = 0

    for inputs, labels in trainloader:

        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(inputs)

        loss = criterion(outputs, labels)

        loss.backward()

        optimizer.step()

        running_loss += loss.item()

        _, predicted = torch.max(outputs, 1)

        total += labels.size(0)

        correct += (predicted == labels).sum().item()

    scheduler.step()

    train_accuracy = 100 * correct / total
```

```
    print(f'Epoch {epoch+1}/{num_epochs}, Loss:
{running_loss/len(trainloader):.4f}, Accuracy: {train_accuracy:.2f}%')

# Testing loop

model.eval()

correct = 0

total = 0

with torch.no_grad():

    for inputs, labels in testloader:

        inputs, labels = inputs.to(device), labels.to(device)

        outputs = model(inputs)

        _, predicted = torch.max(outputs, 1)

        total += labels.size(0)

        correct += (predicted == labels).sum().item()

test_accuracy = 100 * correct / total

print(f'Test Accuracy: {test_accuracy:.2f}%')
```

Output:

Epoch 1/30, Loss: 1.7033, Accuracy: 37.45%

Epoch 2/30, Loss: 1.2892, Accuracy: 53.03%

Epoch 3/30, Loss: 1.1124, Accuracy: 60.24%

Epoch 4/30, Loss: 0.9813, Accuracy: 65.28%

Epoch 5/30, Loss: 0.9020, Accuracy: 68.36%

Epoch 6/30, Loss: 0.8505, Accuracy: 70.40%

Epoch 7/30, Loss: 0.7947, Accuracy: 72.47%

Epoch 8/30, Loss: 0.7466, Accuracy: 74.28%

Epoch 9/30, Loss: 0.7229, Accuracy: 74.92%

Epoch 10/30, Loss: 0.6923, Accuracy: 76.23%

Epoch 11/30, Loss: 0.6657, Accuracy: 77.07%

Epoch 12/30, Loss: 0.6430, Accuracy: 77.82%

Epoch 13/30, Loss: 0.6194, Accuracy: 78.92%

Epoch 14/30, Loss: 0.5912, Accuracy: 79.73%

Epoch 15/30, Loss: 0.5766, Accuracy: 80.24%

Epoch 16/30, Loss: 0.5542, Accuracy: 80.83%

Epoch 17/30, Loss: 0.5422, Accuracy: 81.40%

Epoch 18/30, Loss: 0.5275, Accuracy: 81.86%

Epoch 19/30, Loss: 0.5108, Accuracy: 82.41%

Epoch 20/30, Loss: 0.4951, Accuracy: 82.89%

Epoch 21/30, Loss: 0.4828, Accuracy: 83.44%

Epoch 22/30, Loss: 0.4741, Accuracy: 83.75%

Epoch 23/30, Loss: 0.4590, Accuracy: 84.10%

Epoch 24/30, Loss: 0.4429, Accuracy: 84.82%

Epoch 25/30, Loss: 0.4362, Accuracy: 85.04%

Epoch 26/30, Loss: 0.4242, Accuracy: 85.51%

Epoch 27/30, Loss: 0.4175, Accuracy: 85.54%

Epoch 28/30, Loss: 0.4048, Accuracy: 86.13%

Epoch 29/30, Loss: 0.3925, Accuracy: 86.54%

Epoch 30/30, Loss: 0.3923, Accuracy: 86.40%

Test Accuracy: 86.29%